

radley yeldar.

Cache Me Outside!

Anthony Dang - Technical Director

@anthonydotnet

```
public function getCachedChartBy($get_by, $get_by_value, $date = null, $limit = 100){  
    $cachekey = md5("{$get_by}{$get_by_value}{$limit}{$date}");  
    return Cache::remember($cachekey, rand(30,300), function() use ($get_by, $get_by_value, $date, $limit){  
        return $this->getChartBy($get_by, $get_by_value, $date, $limit);  
    });  
}
```

radley yeldar.

A bit about me



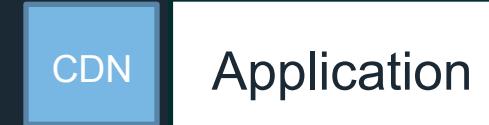
UMBRACO MVP

- Limited bandwidth / availability
 - Cost
 - Blazing fast UX
 - SEO
 - Poor Performance?
-

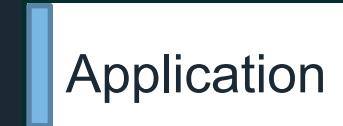
DNS/Proxy



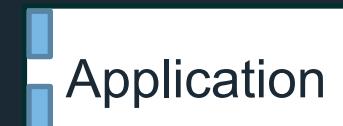
Network



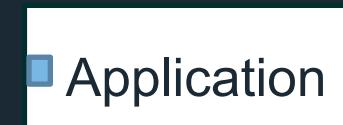
Page Output



Donut Caching



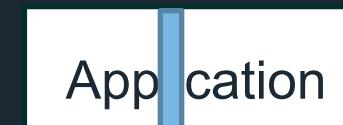
Donut Hole Caching



Methods



Service layer



Data layer

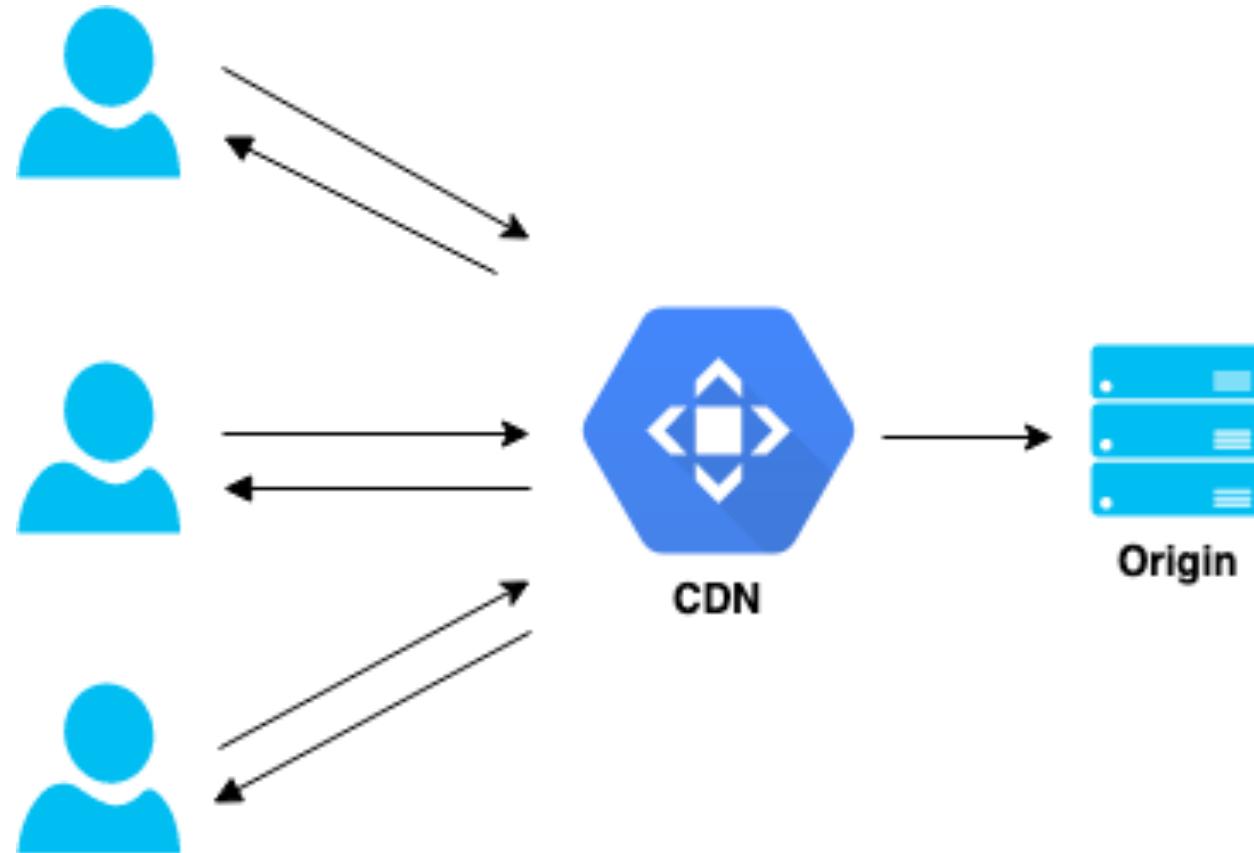


Pre-Indexing



Multi-layered



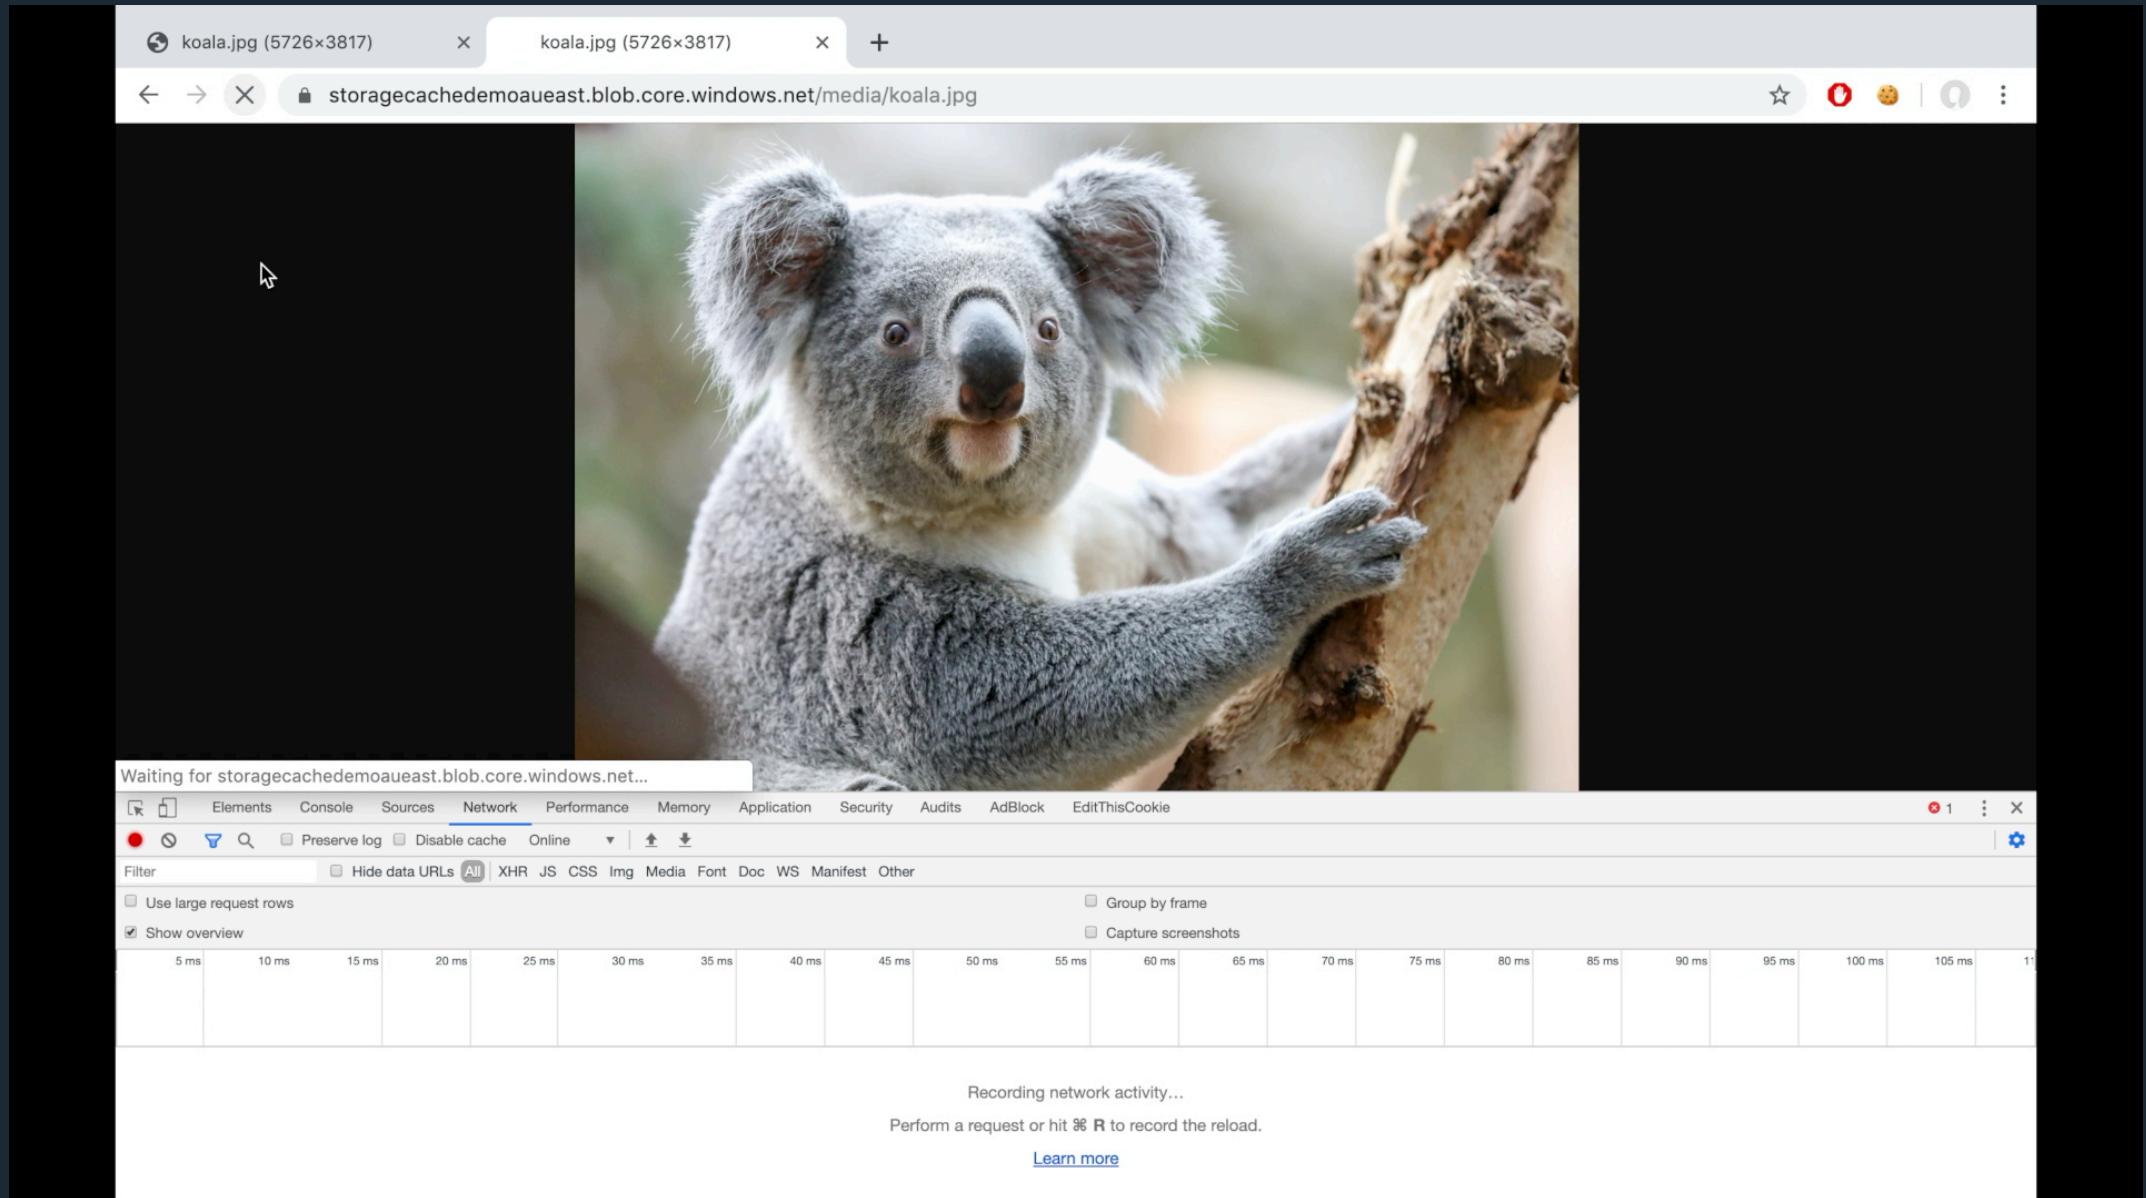


The screenshot shows the Microsoft Azure portal interface for managing Content Delivery Network (CDN) profiles. The left sidebar contains a navigation menu with various service icons, including 'Create a resource', 'Home', 'Dashboard', 'All services', 'FAVORITES' (with 'CDN profiles' checked), 'All resources', 'Resource groups', 'App Services', 'Function App', 'Virtual machines', 'SQL databases', 'Azure Cosmos DB', 'Storage accounts', 'CDN profiles' (which is the current selected item), 'Service Bus', 'Virtual networks', 'Load balancers', 'Azure Active Directory', 'Monitor', and 'Advisor'. The main content area has a title 'CDN profiles' under 'Default Directory'. It features a search bar and filter options for 'Subscriptions: Pay-As-You-Go' (Filter by name, All resource groups, All locations, All tags, No grouping). Below this, it displays '0 items' and a table header with columns: NAME, PRICING TIER, RESOURCE GROUP, LOCATION, and SUBSCRIPTION. In the center, there is a large cloud icon and the message 'No cdn profiles to display'. A link 'Try changing your filters if you don't see what you're looking for.' is provided, along with a blue 'Create cdn profile' button.

Blob Storage vs Content Delivery Networks

radley yeldar.

8

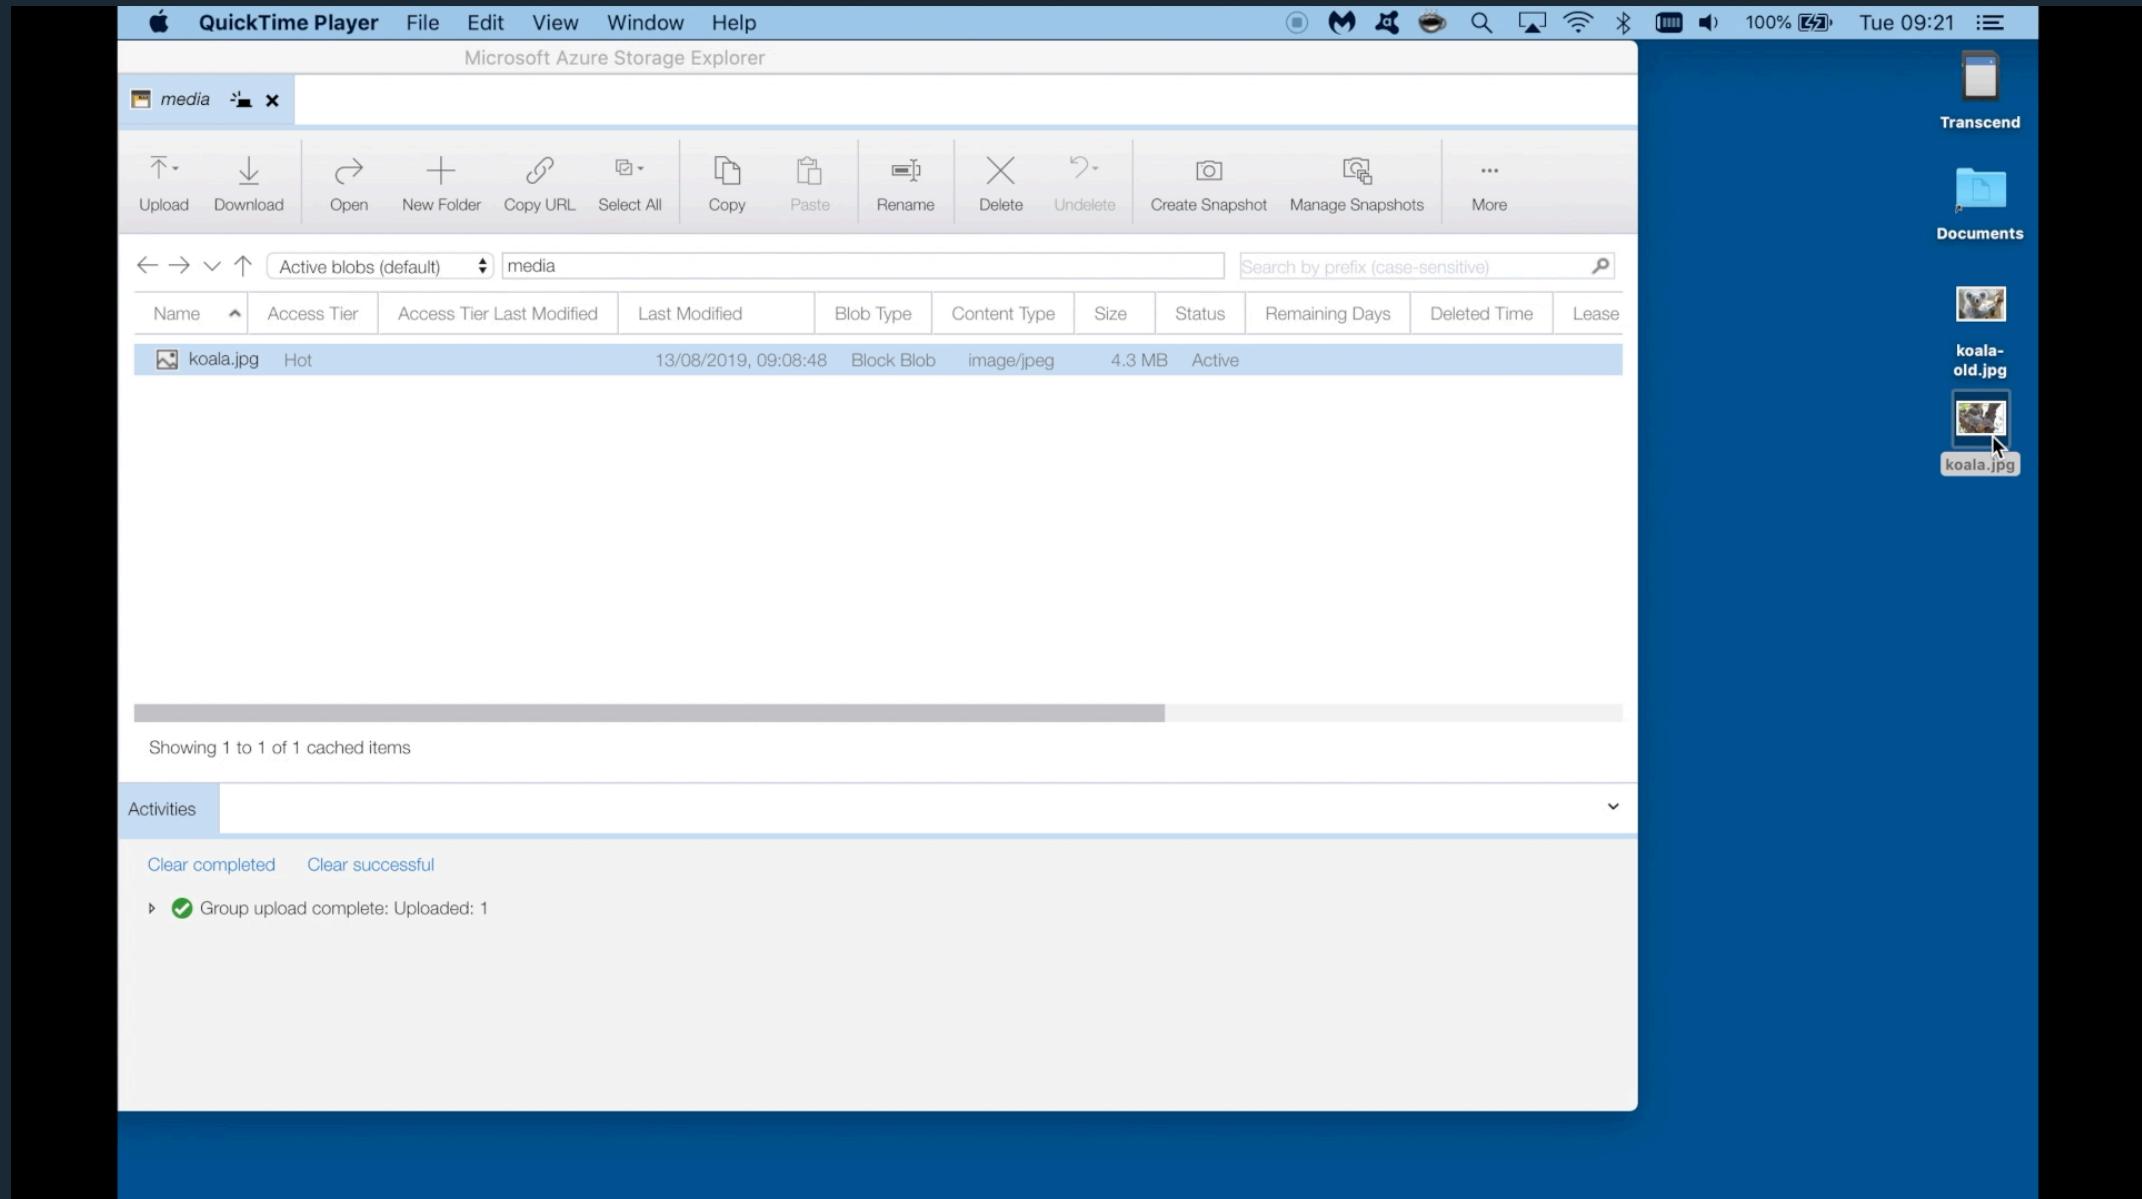


@anthonydotnet

Content Delivery Networks - Cache Update & Purging

radley yeldar.

9

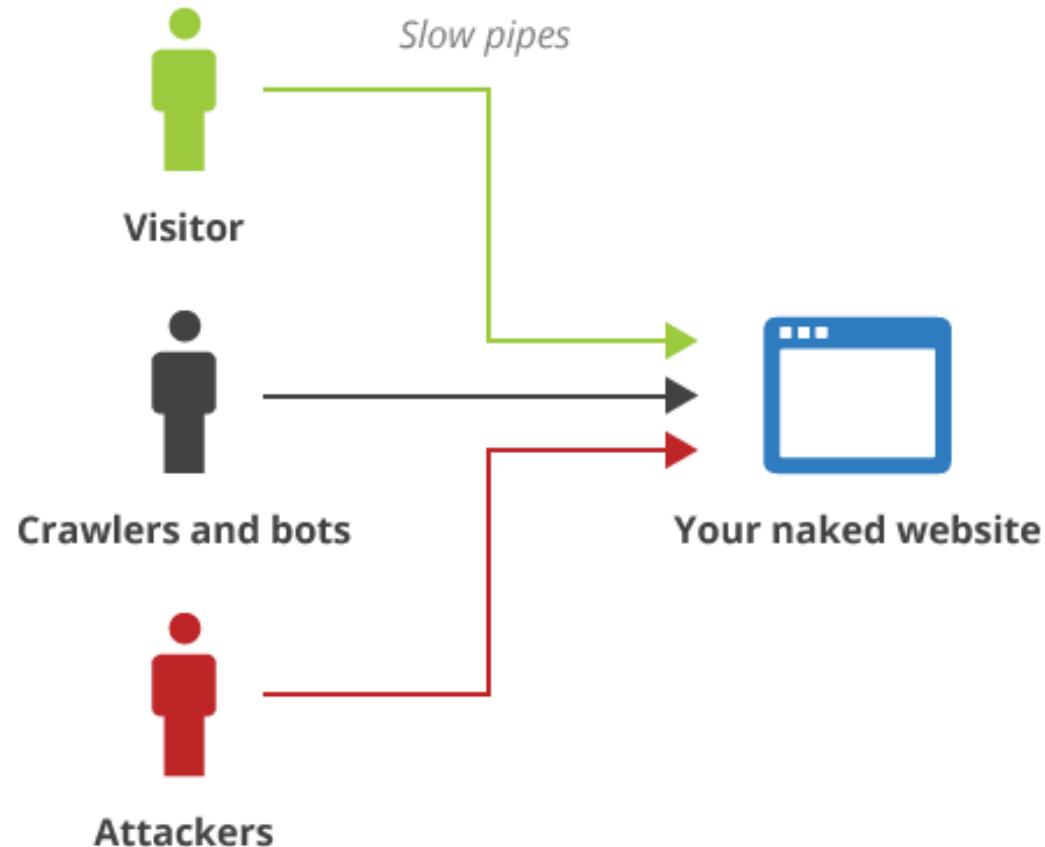


- When will cache timeout?
- Is stale content acceptable?
- How to clear the cache?
- How to get files up there?

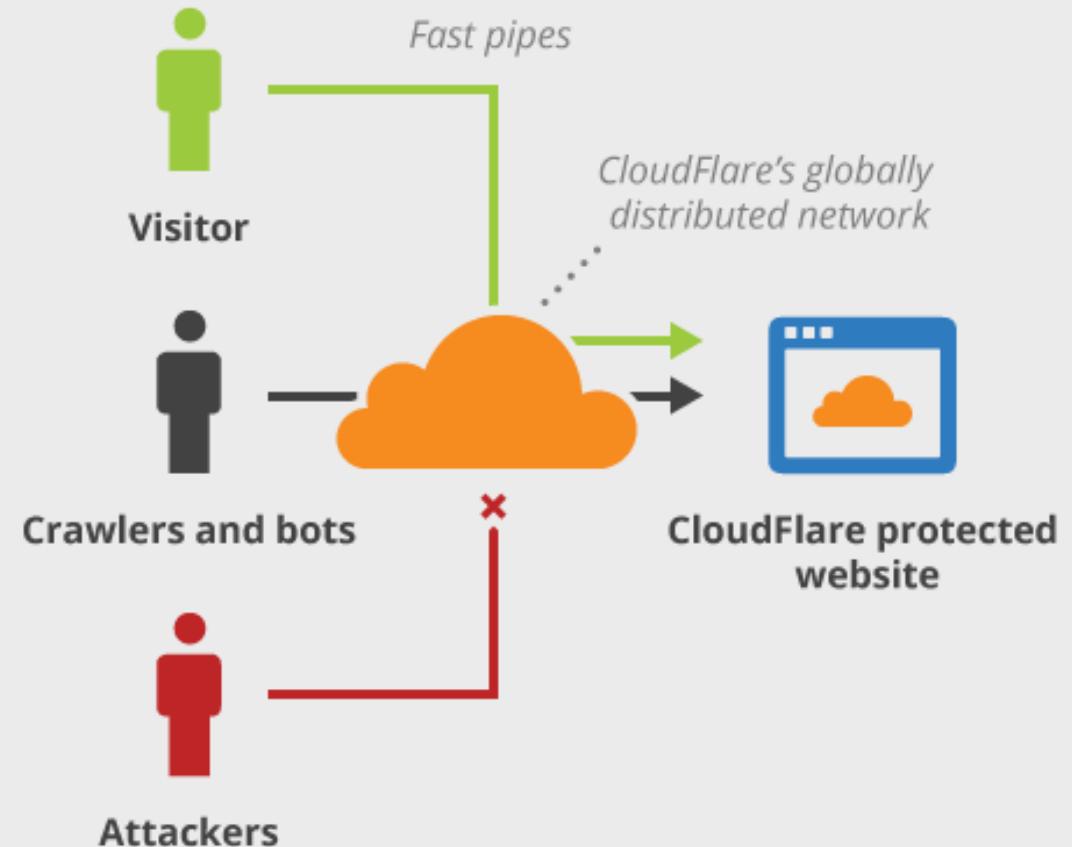
Browser cache?

- Is this a problem?

Without CloudFlare



With CloudFlare



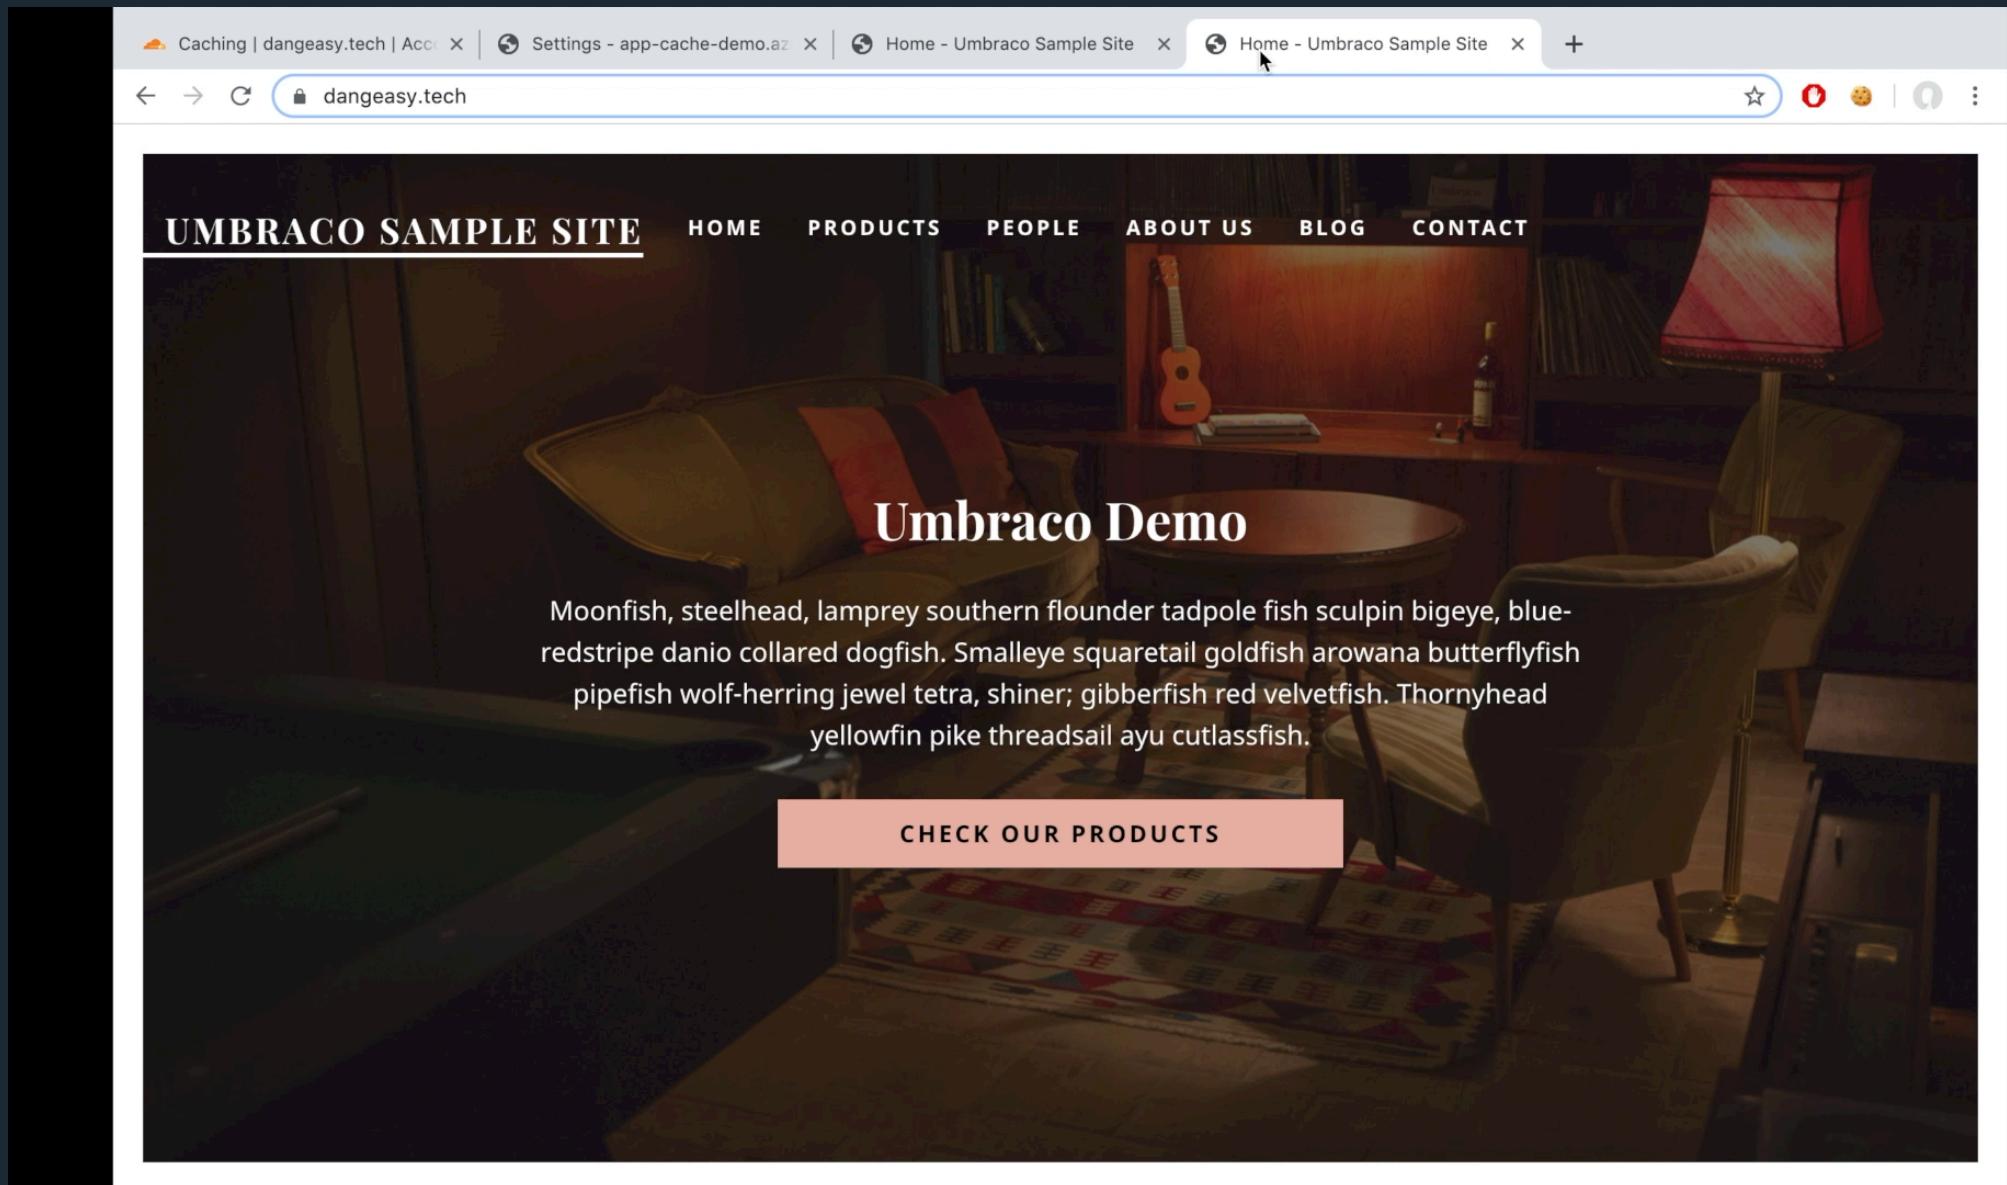
The screenshot shows the Cloudflare dashboard for the site `dangeeasy.tech`. At the top, there's a navigation bar with icons for Overview, Analytics, DNS, Crypto, Firewall, Access, Speed, Caching, Workers, Page Rules, Network, Traffic, Stream, Custom P..., Apps, and Scrape S... The main area displays real-time statistics:

- Unique Visitors:** 8
- Total Requests:** 9
- Percent Cached:** 0.00%
- Total Data Served:** 26 kB
- Data Cached:** 0 B

Below these stats is a chart showing traffic over the period from 13 AUGUST — 14 AUGUST. The chart shows a peak in visitors and requests around August 14th. The dashboard also features a "Notifications" section with two items:

- Speed page**: A new feature notification encouraging users to try the Speed page. It includes a "Dismiss" button.
- Billing update**: A message about upgrading the billing system, with a "common questions" link and a "Dismiss" button.

In the bottom right corner, there are "Quick Actions" buttons for "Purge Cache" and "DNS Settings". A "Under Attack Mode" status is shown with a "Send Feedback" button. The footer of the dashboard includes the Cloudflare logo and the site name `dangeeasy.tech`.



- Hides downtime/crash until timeout
 - Errors being cached
 - Limited page rules
 - 100, then 1USD per rule
 - Security
-

- Self hosted – Linux
- Cached HTTP Proxy
- Granular unlimited rules
- Configuration



The screenshot shows a GitHub repository page for `varnish-4.0-configuration-templates`. The repository was forked from `mattiasgeniar/varnish-4.0-configuration-templates`. The current branch is `master`, and the file displayed is `default.vcl`. The commit `c59b06b` was made on 22 Oct 2018 by `hvelarde`, adding `fbclid` to the list of parameters to be removed from URLs. The code editor shows 414 lines of VCL code, which includes defining a backend, setting host and port, and defining a probe for the backend.

```
1 vcl 4.0;
2 # Based on: https://github.com/mattiasgeniar/varnish-4.0-configuration-templates/blob/master/default.vcl
3
4 import std;
5 import directors;
6
7 backend server1 { # Define one backend
8     .host = "127.0.0.1";      # IP or Hostname of backend
9     .port = "80";              # Port Apache or whatever is listening
10    .max_connections = 300; # That's it
11
12    .probe = {
13        #.url = "/"; # short easy way (GET /)
14        # We prefer to only do a HEAD /
15        .request =
```

- Crazy config file
 - Maintain linux server
 - Fastly?
-

Coding time

```
@Html.CachedPartial("MyPartialName", new MyModel(), 3600)
```

- Hardcoded cache time
 - How to clear cache?
 - How to turn off cache?
-

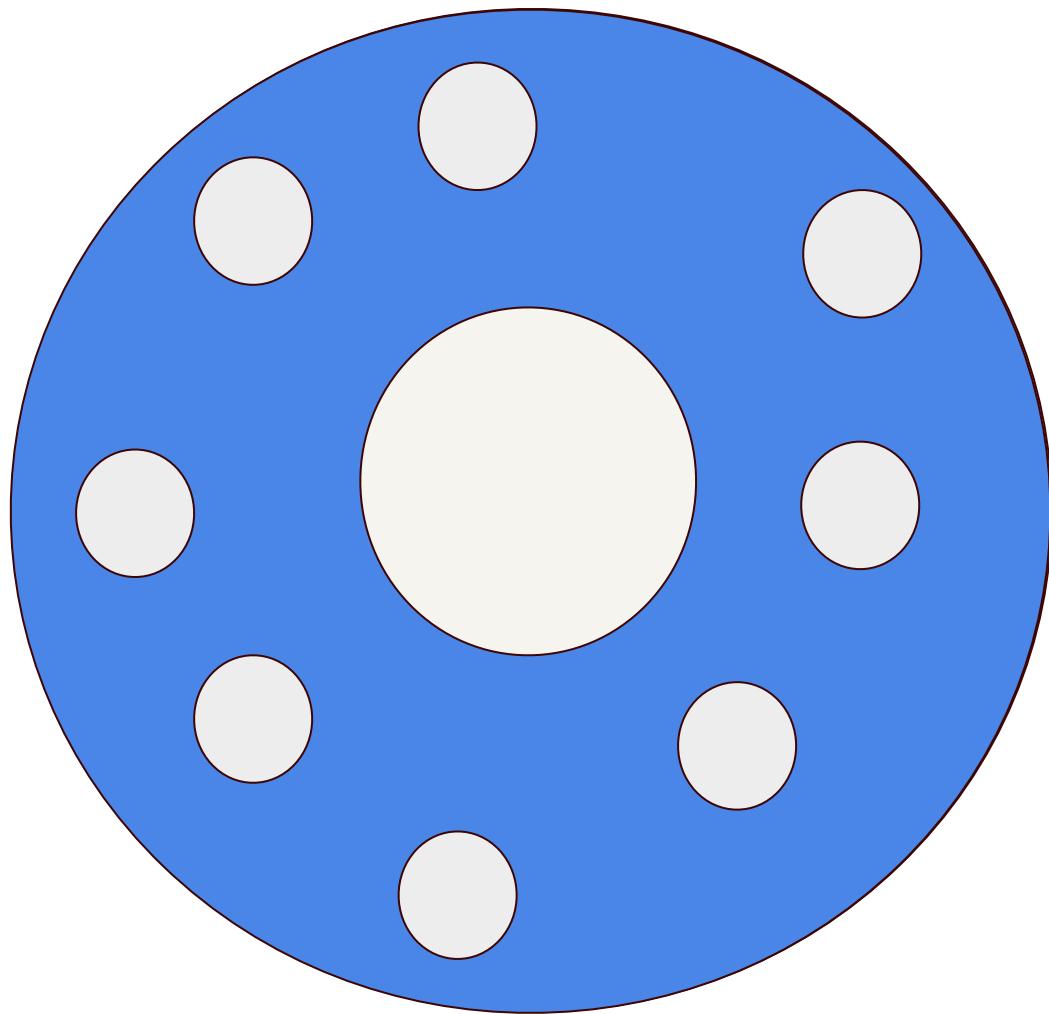
```
@Html.CachedPartial("MyPartialName", new MyModel(), 3600)
```

```
public class ProfileController : Controller
{
    [OutputCache(CacheProfile="Cache1Hour")]
    public string Index()
    {
        return DateTime.Now.ToString("T");
    }
}

<caching>
    <outputCache enableOutputCache="true" />
    <outputCacheSettings>
        <outputCacheProfiles>
            <add name="Cache1Hour" duration="3600" varyByParam="none" />
        </outputCacheProfiles>
    </outputCacheSettings>
</caching>
```

- Remember to cache by param etc.
- Caches everything!
- Cache not close to the problem?
- Harder to diagnose performance problems

```
public class ProfileController : Controller
{
    [OutputCache(CacheProfile="Cache1Hour")]
    public string Index()
    {
        return DateTime.Now.ToString("T");
    }
}
```



```
@Html.Action("Login", "Account", true)
```

```
[DonutOutputCache(Duration=60)]  
public ActionResult Index()
```

excludeFromParentCache

```
[DonutOutputCache(CacheProfile="TwoMins")]  
public ActionResult Index()
```

```
[DonutOutputCache(Duration=60, VaryByCustom="whatever")]  
public ActionResult Index()
```

```
[DonutOutputCache(Duration=60, VaryByParam="something;that")]  
public ActionResult Index(string something)
```



- Remember to punch holes
- Cache not close to the problem?
- Harder to diagnose performance problems

```
[DonutOutputCache(Duration=60)]
```

```
public ActionResult Index()
```

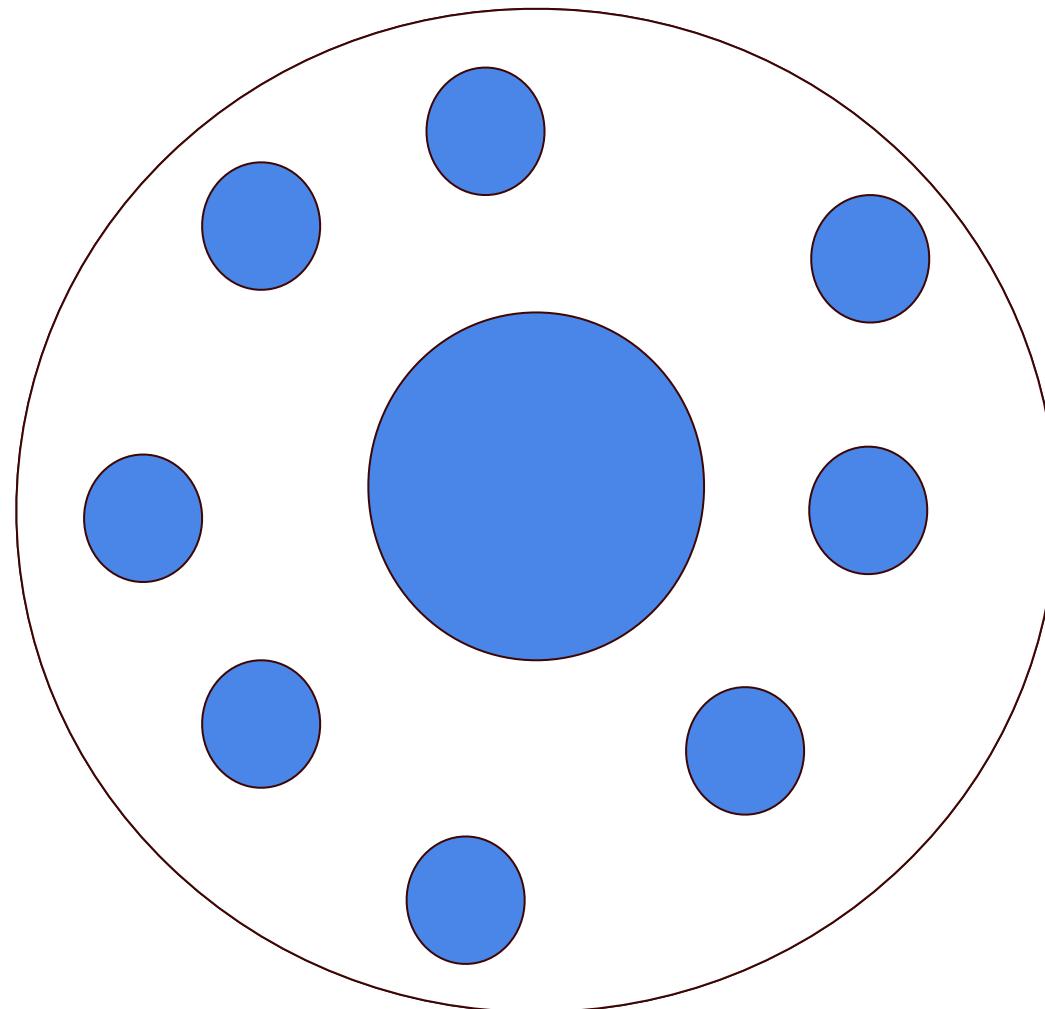
```
[DonutOutputCache(CacheProfile="TwoMins")]
```

```
public ActionResult Index()
```

```
[DonutOutputCache(Duration=60, VaryByCustom="whatever")]
```

```
public ActionResult Index()
```





```
<div>@Html.Action("CategoriesList")</div>
```

```
[ChildActionOnly]
[OutputCache(Duration=60)]
public ActionResult CategoriesList()
{
    // Get categories list from the database and
    // pass it to the child view
    ViewBag.Categories = Model.GetCategories();
    return View();
}
```

- Cache not close to the problem?
- Harder to diagnose performance problems
- No granular control

```
[ChildActionOnly]
[OutputCache(Duration=60)]
public ActionResult CategoriesList()
{
    // Get categories list from the database and
    // pass it to the child view
    ViewBag.Categories = Model.GetCategories();
}
```

```
public object GetValue(string id)
{
    var value = GetSomethingExpensive(id);

    return value;
}
```

```
public object GetValue(string id)
{
    var cacheKey = id;

    object value = _cache.Get<object>(cacheKey);

    if (value == null)
    {
        value = GetSomethingExpensive(id);
        _cache.Add(cacheKey, value, new TimeSpan(ONE_HOUR));
    }

    return value;
}
```

- Not single responsibility
- Repeated code
- Where is cache used?
- How to clear / turn off?
- Diagnose performance?



```
public object GetValue(string id)
{
    var cacheKey = id;

    object value = _cache.Get<object>(cacheKey);

    if (value == null)
    {
        value = GetSomethingExpensive(id);
        _cache.Add(cacheKey, value, new TimeSpan(0, 0, 5));
    }

    return value;
}
```

- Dependency Injection & Interfaces
- Service Layer
- Cached proxy
- Replace concrete implementation at startup

```
public interface IExampleService
{
    object GetValue(string id);
}
```

```
public class ExampleService : IExampleService
public class ExampleServiceCachedProxy : IExampleService
```

```
private static void RegisterCachedServices(Composition composition)
{
    if (ConfigurationHelper.ServiceCacheEnabled())
    {
        composition.Register(typeof(ExampleService), typeof(ExampleService));

        composition.Register(typeof(ICache), typeof(Cache));
        composition.Register(typeof(IExampleService), typeof(ExampleServiceCachedProxy));
    }
    else
    {
        composition.Register(typeof(IExampleService), typeof(ExampleService));
    }
}

<add key="ServiceCache:Enabled" value="true" />
```

```
public object GetValue(string id)
{
    var value = GetSomethingExpensive(id);

    return value;
}
```

```
public object GetValue(string id)
{
    var cacheKey = $"{typeof(ExampleServiceCachedProxy)}_{id}";

    return _cache.Get(cacheKey, () => _exampleService.GetValue(id));
}
```

```
public void Initialize()
{
    ContentService.Published += ContentService_Published;
}

private void ContentService_Published(IContentService sender,
    ContentPublishedEventArgs e)
{
    _cache.RemoveByPrefix(typeof(ExampleServiceCachedProxy).ToString());
}
```

- Learning curve can be hard
 - IoC init get be bloated
 - Some methods only pass through
 - No way to manually clear
-

```
internal class DefaultRepositoryCachePolicy<TEntity, TId> : RepositoryCachePolicyBase<TEntity, TId>
|   where TEntity : class, IEntity

public override TEntity Get(TId id, Func<TId, TEntity> performGet, Func<TId[], IEnumerable<TEntity>> performGetAll)
{
    var cacheKey = GetEntityCacheKey(id);
    var fromCache = Cache.GetCacheItem<TEntity>(cacheKey);

    // if found in cache then return else fetch and cache
    if (fromCache != null)
        return fromCache;
    var entity = performGet(id);

    if (entity != null && entity.HasIdentity)
        InsertEntity(cacheKey, entity);

    return entity;
}
```

- All previous problems
- Stale entities – Concurrent write & reads?
- Difficult to distribute (e.g. multi web-heads)

Consider your architecture first!

- External Key-Value storage service
- Great alternative to local memory
- Great for large amounts of data



redis

★ DNS name

redis-cachedemo



.redis.cache.windows.net

★ Subscription

Pay-As-You-Go

★ Resource group i

rg-app-cache-demo

[Create new](#)

★ Location

UK South

★ Pricing tier ([View full pricing details](#))

Basic C0 (250 MB Cache)



(PREVIEW) Availability zone

Requires Premium tier

Redis Cluster i>
Requires Premium tierData persistence i>
Requires Premium tierVirtual Network i>
Requires Premium tier Unblock port 6379 (not SSL encrypted)[Create](#)[Automation options](#)



CACHE NAME	CACHE SIZE	NETWORK PERFORMANCE	NUMBER OF CLIENT CONNECTIONS	PRICE
C0	250 MB	Low	256	~£14.963/month
C1	1 GB	Low	1,000	~£37.542/month
C2	2.5 GB	Moderate	2,000	~£60.937/month
C3	6 GB	Moderate	5,000	~£122.418/month
C4	13 GB	Moderate	10,000	~£142.548/month
C5	26 GB	High	15,000	~£285.640/month
P1	6 GB	Moderate	7,500	~£150.709/month
P2	13 GB	High	15,000	~£301.962/month
P3	26 GB	High	30,000	~£603.380/month
P4	53 GB	Highest	40,000	~£1,207.848/month

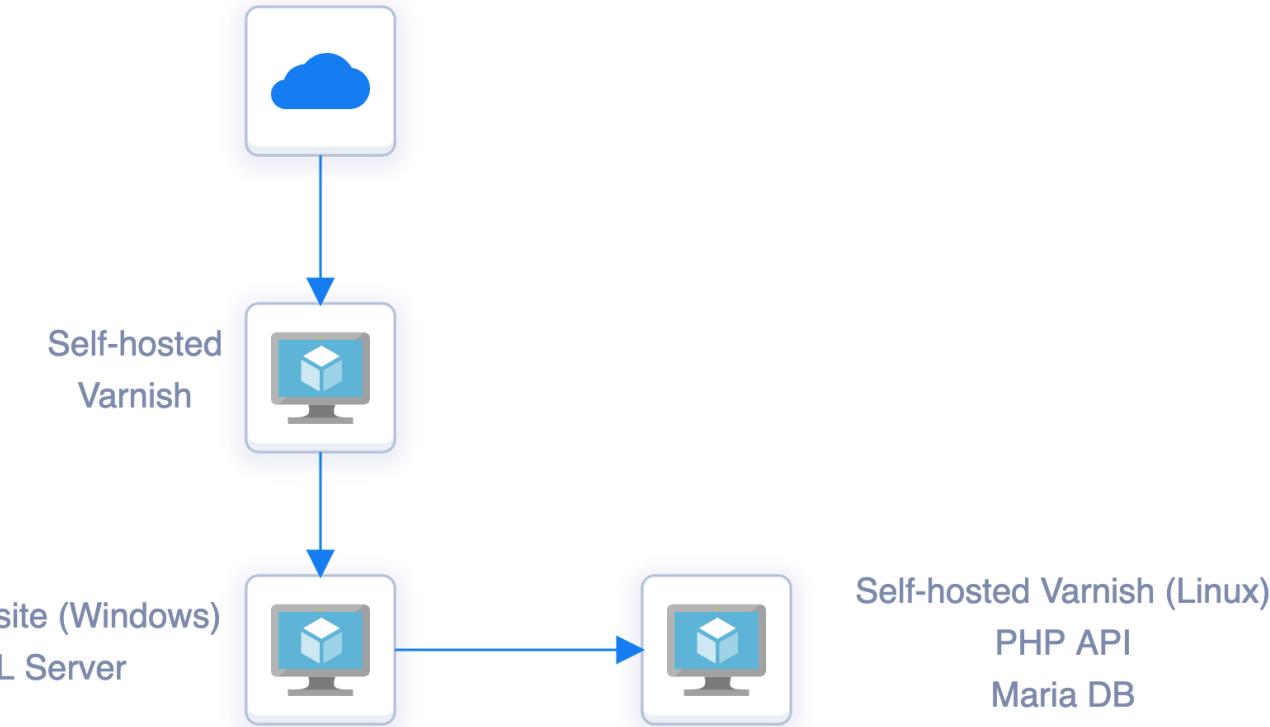
- Cost
 - More infrastructure
 - Uptime
-



redis



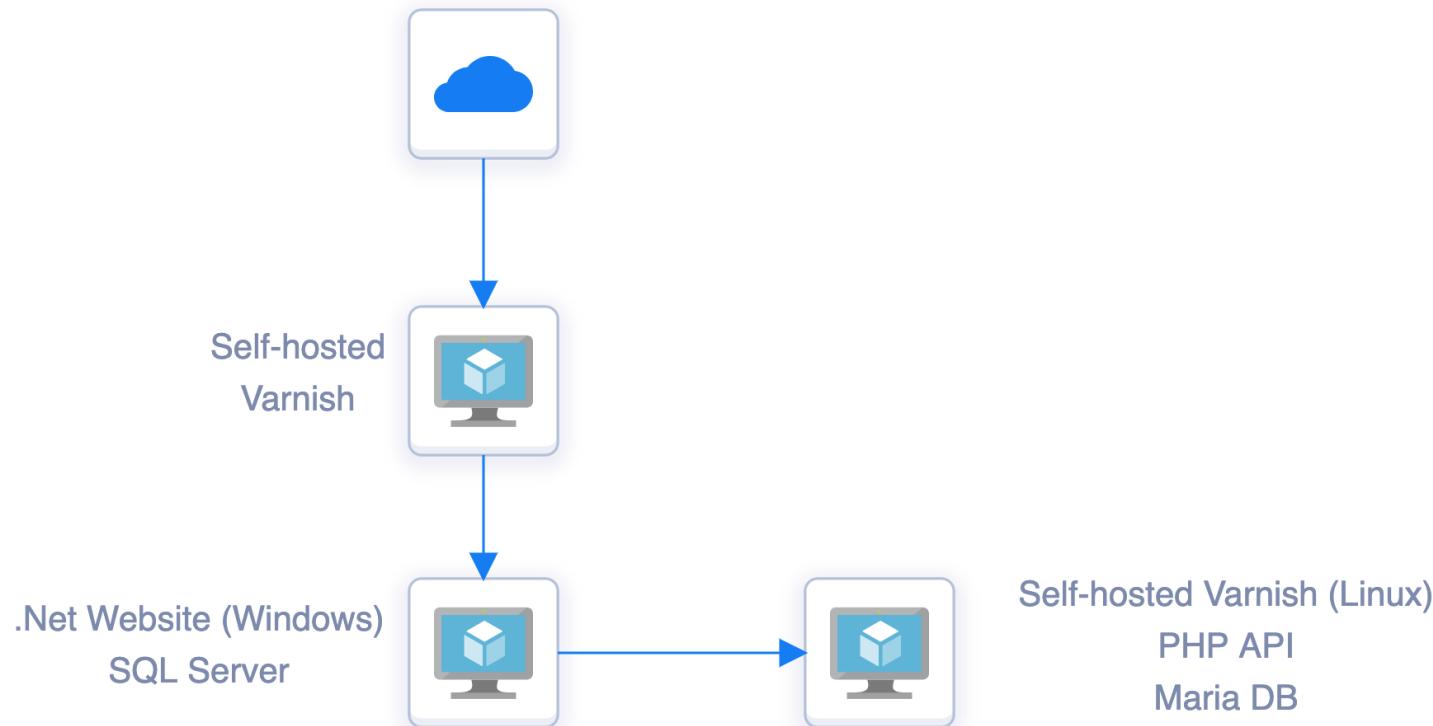
Varnish (VM 1)	24hr
Website: Memory Cache (VM 2)	1hr
Varnish (VM 3)	1hr
API: Random Memory Cache (VM 3)	10min



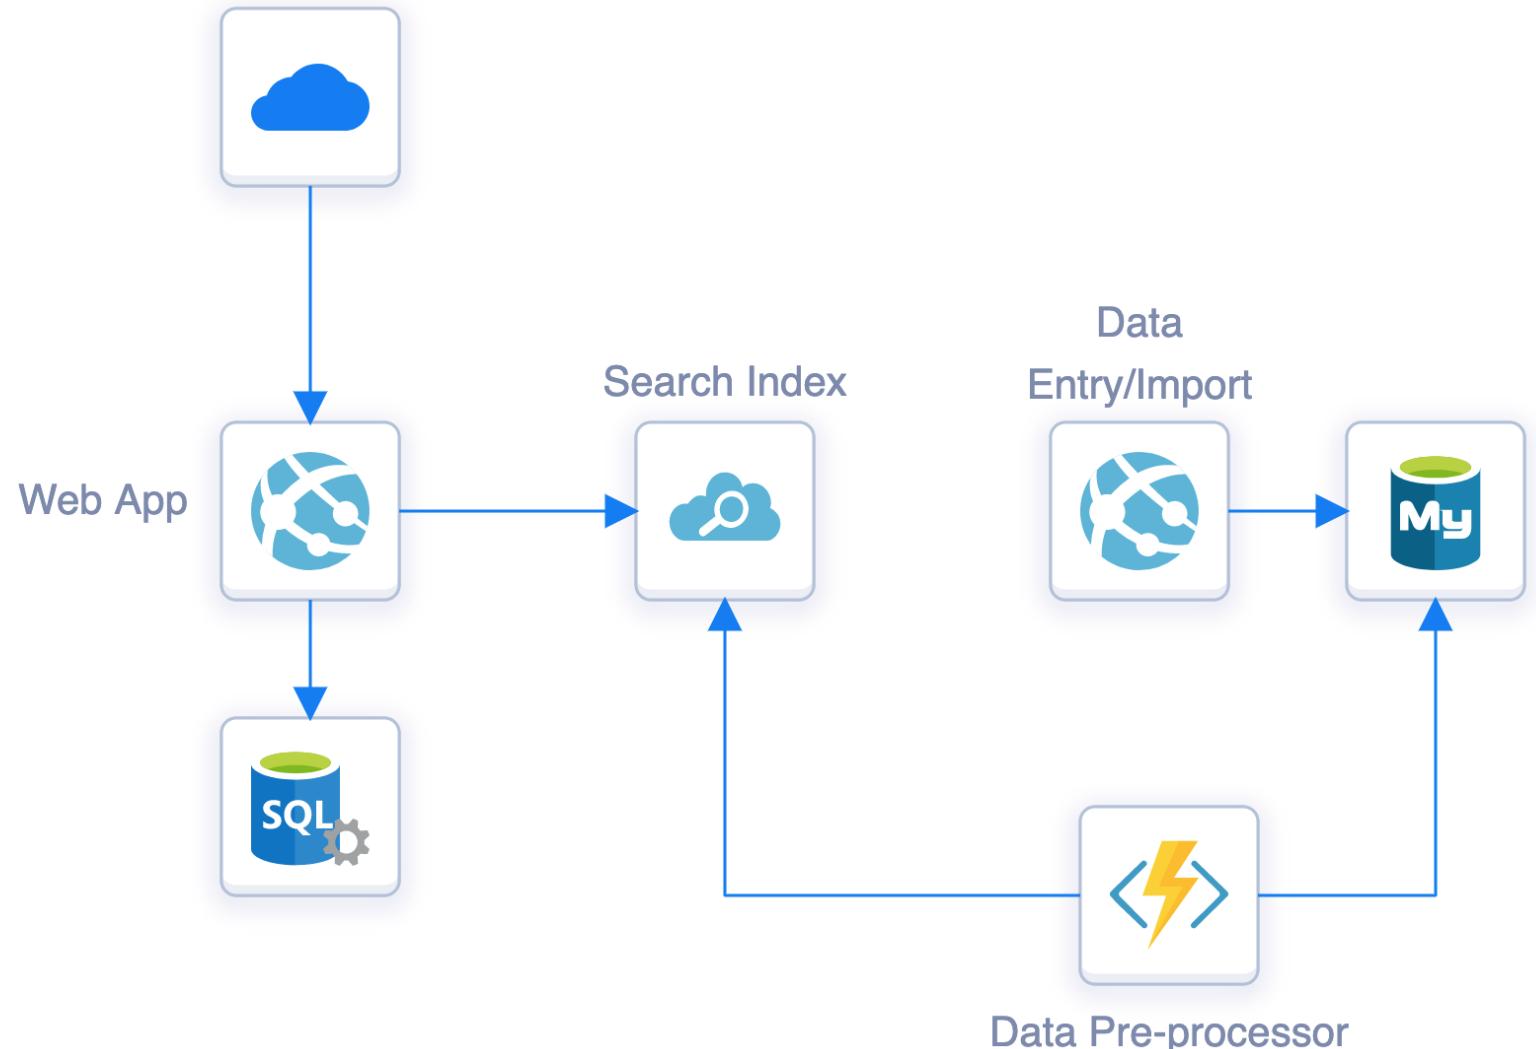
Considerations

- What is the actual cache time?
- How to clear all the cache?

- What was the problem?
- Complex / inefficient SQL Queries



- Pre-process
- Index



More infrastructure => More downtime

1 Service (99.95%)

21m 54.9s per month

2 Services (99.95% each)

$99.95 \times 99.95 = 99.90\%$: 43m 49.7s

3 Services (99.95% each)

$99.95 \times 99.95 \times 99.95 = 99.85\%$: 1h 5m 44.6s

- Can add instability/downtime
- Clearing multiple layers simultaneously
- Pre-index up to date?

Cache is not the answer to everything!!

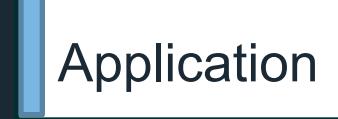
DNS/Proxy



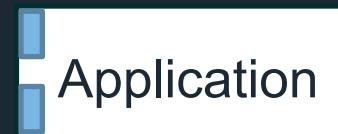
Network



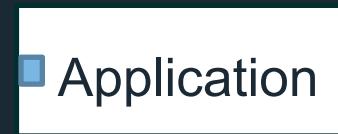
Page Output



Donut Caching



Donut Hole Caching



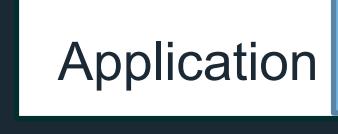
Methods



Service layer



Data layer



Pre-Indexing



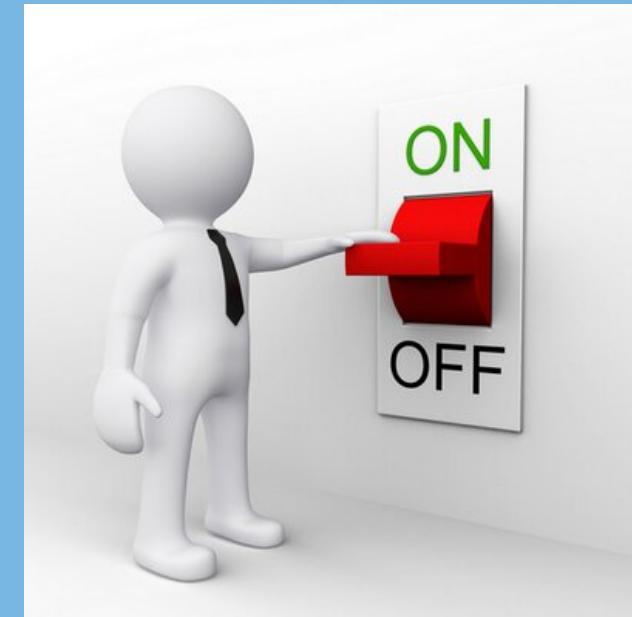
Multi-layered



- Cache as a last resort
- Favour pre-processing & indexing
- Automatic & manual clearing
- Avoid caching everything
- Cache close to the problem
- Keep it simple and obvious

Multiple layers => multiple headaches

No perfect solution, only trade-offs



Thanks!

radley yeldar.

Cache me on Twitter
@anthonydotnet

Git Kraken: anthonyd19

@anthonydotnet